

What does the language actually cost?

A conformance-verified benchmark of ten TypeScript, Go and Rust implementations of one database-bound business transaction

Halleluyah Oludele

22 September 2026

Abstract

Comparisons between server runtimes usually measure microbenchmarks — JSON serialisation, a single-row fetch — and rarely establish that the implementations being compared do the same work. We specify one composite business transaction (order pricing and checkout: six reads, a stacking discount engine, and an eleven-statement write transaction), implement it ten times across TypeScript, Go and Rust, and verify byte-for-byte output equivalence across all ten before any timing is recorded. Measuring 450 runs under pinned CPU sets with a coordinated-omission-safe load generator, we find that at a fixed 300 requests per second the ten stacks span 1.0 ms of p99 latency, and that Rust and Go are statistically indistinguishable ($0.99\times$, $p=0.71$). The separations that do survive correction are between language tiers and, larger than any of them, between data-access strategies: an ORM costs $1.84\times$ throughput against the identical router over hand-written SQL. A zero-dependency Bun service outperforms both third-party TypeScript stacks. Code, data and analysis are public.

1 What this measures, and why

Choosing a backend language is usually argued with throughput numbers drawn from benchmarks that do very little work per request. The TechEmpower Framework Benchmarks [1], the reference point for this comparison, cover 331 frameworks; their database tests fetch one row, or twenty rows, from a single table. A service that spends most of its time on database round-trips and application-side business logic is a different regime, and the ranking from one does not transfer to the other.

Two further gaps motivated this work. First, implementations in a multi-language benchmark are usually accepted on review rather than proven equivalent, so a stack can place well by doing marginally less work. Second, results are typically published as point estimates, with no repetition variance and no significance testing, which makes it impossible to tell an ordering from noise.

This is a systems and artifact contribution, not a novel finding: that compiled languages outperform interpreted ones under saturation is not news. What the artifact provides is a workload shaped like real business code, a conformance gate that makes the comparison meaningful, and an analysis that says which differences are real.

2 Workload

The contract is specified once, in `spec/SPEC.md`, and implemented to the letter by all ten services. Three endpoints:

Pricing (`POST /api/quote`) issues six reads — customer joined to tier, products by SKU, active discount rules, regional tax rates, an optional coupon, and inventory — then runs a pure-CPU

engine. Each cart line is priced against every active rule; stacking resolves to the largest non-stackable rule plus all stackable ones, capped at 60% of the line; a coupon is spread across lines pro rata by net value with remainder cents assigned to the largest line; per-line tax is recomputed on the reduced amount; shipping is banded by total weight; loyalty points accrue by tier multiplier. All money is integer minor units and every basis-point multiplication rounds half-up on the final cent, so the result is exactly reproducible across languages.

Checkout (`POST /api/checkout`) runs the same engine inside one transaction, taking `SELECT ...FOR UPDATE` on inventory in product-id order to avoid deadlock, rejecting insufficient stock with HTTP 409, then writing the order, a multi-row line insert, per-line inventory reservations, two ledger entries, a coupon usage increment and a loyalty transaction before committing.

Summary (`GET /api/customers/:id/summary`) performs four reads over a customer’s last 50 orders and their lines and aggregates in process into lifetime value, top categories by spend, points balance and a segment label.

The dataset holds 50,000 customers, 5,000 products, 200,000 historical orders and 600,247 order lines, generated deterministically from a fixed seed.

3 Implementations

Table 1: The ten stacks. Within a language the pricing engine is a single shared module, so framework rows differ only by framework.

id	Language	Framework	Driver
ts-bun-only	TypeScript	hand-rolled router	Bun.sql (built in, zero deps)
ts-bun-hono	TypeScript	Hono	pg (npm)
ts-node-fastify	TypeScript	Fastify	pg (npm)
go-nethttp	Go	net/http (stdlib)	pgx/v5
go-fiber	Go	Fiber (fasthttp)	pgx/v5
go-gin	Go	Gin	pgx/v5
go-gin-gorm	Go	Gin	GORM
rust-hyper	Rust	hyper (none)	tokio-postgres
rust-axum	Rust	Axum	tokio-postgres
rust-actix	Rust	Actix Web	tokio-postgres

Two rows are deliberate exceptions to the shared-module rule. `go-gin-gorm` reaches the same logic through an ORM, isolating what the ORM costs when the router is held constant. `ts-bun-only` has no third-party dependency at all: `Bun.serve` for HTTP, `Bun.sql` for Postgres, and a sixty-line router. It runs from a bare directory with no `node_modules` and no lockfile.

4 Conformance

Before any measurement, `bench/verify.py` replays a fixed corpus against each service in turn and diffs the canonical JSON responses across all ten. A stack that skipped a query, rounded differently, or computed a different invoice fails here. Two normalisations are applied: the order id, which is a database sequence and differs per run, and ISO timestamp spelling, since `:00Z` and `:00.000Z` denote the same instant and the specification does not dictate fractional-second formatting. All ten implementations pass.

5 Method

Hardware and isolation. Intel Core i7-8665U (4 cores, 8 threads), 31 GB RAM, Arch Linux 7.1.9. The three parties never share a core: PostgreSQL 18 is confined by Docker to cores 0–2, the service under test by `taskset` to 3–5, and the load generator to 6–7. The CPU governor is pinned to `performance`; under the default `powersave` governor an earlier iteration of this suite showed a 66% median spread across repetitions.

Database. PostgreSQL 18 in Docker with fixed settings (`shared_buffers=4GB`, `synchronous_commit=off`, `max_connections=400`). Disabling synchronous commit is a measurement choice: it removes SSD commit stalls that would otherwise dominate the write workload. Autovacuum is disabled on the six churned tables and vacuum runs explicitly between runs; autovacuum firing mid-measurement was the largest single source of variance we identified. The database is restored to seeded state before every run.

Load generation. A purpose-built generator runs closed-loop (64 connections, send as fast as the server answers) to find the saturation ceiling, and open-loop at a fixed arrival rate to measure latency under controlled load. In open-loop mode latency is measured from each request’s *intended* send time rather than its actual dispatch, which is the standard correction for coordinated omission: a server that falls behind cannot hide the delay inside its own slowness. Warmup is excluded from statistics so JIT-compiled runtimes are not charged for their first requests.

Design. 5 repetitions of 10 s after 3 s warmup, across three workloads and three load profiles, for 450 runs in total. Run order is shuffled within each repetition so thermal drift cannot systematically favour whichever stack ran first. Every run records a machine-speed probe — fixed CPU work, timed — so runs taken while the machine was busy remain identifiable; 58 of 450 runs were flagged, and excluding them moves the reported medians by 0.0% at the median.

6 Statistical treatment

With ten treatments measured over nine tasks (three workloads $\times$ three load profiles), the appropriate omnibus test is the Friedman test with Nemenyi post-hoc [2], whose power derives from the number of tasks rather than the number of repetitions, and which assumes no error distribution. We report medians and bootstrap intervals for transparency but base significance claims on the rank-based tests.

Friedman rejects the null for p99 latency ($\chi^2=56.7$, $p=5.8\text{e-}09$), for CPU per request ($\chi^2=77.4$, $p=5.2\text{e-}13$) and for peak memory ($\chi^2=80.6$, $p=1.2\text{e-}13$), each over $N=9$ tasks and $k=10$ treatments. The Nemenyi critical difference at the 5% level is 4.52 ranks.

Five pairwise contrasts were specified before data collection, each corresponding to one design question. For these, throughput is normalised within each workload by that workload’s grand median across stacks and pooled, giving 15 independent runs per stack; they are compared with the Mann–Whitney U test under Holm correction across the family of five. Normalising within workload is what makes pooling legitimate — the three workloads differ by almost an order of magnitude in absolute throughput — and the power comes from blocking, not from re-using observations.

Throughput at saturation – 64 connections, flat out

Median of 5 repetitions; whiskers are the bootstrap 95% interval. Ember marks the zero-dependency Bun service.




Figure 1: Throughput at saturation. Whiskers are bootstrap 95% intervals over 5 repetitions.

7 Results

7.1 Under realistic load, the runtime is not the variable

At a fixed 300 requests per second the ten stacks span 1.0 ms of p99 latency, from 2.64 ms (`go-nethttp`) to 3.60 ms (`ts-bun-hono`). Figure 2 shows that the spread only opens as offered load approaches each stack’s ceiling.

Where each stack starts to hurt

p99 latency as offered load rises. Under a fixed rate the stacks are nearly indistinguishable; the spread only opens at saturation.




Figure 2: p99 latency as offered load rises.

7.2 Rank analysis

7.3 Planned contrasts

Three findings survive correction. **Rust and Go are indistinguishable on this workload:** `rust-axum` against `go-fiber` gives $0.99\times$ (95% CI [0.86, 1.07], $p=0.71$). **The zero-dependency Bun service outperforms both third-party TypeScript stacks**, by $1.33\times$ against Hono with node-postgres and $1.72\times$ against Node with Fastify. **The ORM is the single largest cost on the board:** holding the Gin router constant and changing only the data layer gives $1.84\times$ throughput, and on the write workload GORM consumes 2.87 ms of CPU per request against 0.83 ms.

Table 2: Saturated throughput (req/s) and p99 (ms) per workload; CPU milliseconds per request and peak resident memory are for the pricing workload.

stack	Quote		Checkout		Summary		CPU	RSS
	req/s	p99	req/s	p99	req/s	p99	ms	MB
rust-axum	3,919	21.6	1,366	63.5	5,474	15.7	0.281	8
rust-actix	3,867	21.4	1,382	62.4	5,257	17.0	0.317	9
rust-hyper	3,680	24.3	1,438	59.4	5,268	16.5	0.271	8
go-fiber	3,637	23.1	1,597	51.8	5,313	16.1	0.409	22
go-nethttp	3,546	23.6	1,572	55.7	5,031	17.4	0.444	23
go-gin	3,495	24.8	1,495	62.0	5,127	17.1	0.457	31
go-gin-gorm	2,271	50.8	565	389.1	3,203	40.0	1.125	35
ts-bun-only	2,243	37.0	1,048	83.3	3,699	22.9	0.401	79
ts-bun-hono	1,850	44.0	786	110.8	2,913	29.3	0.518	92
ts-node-fastify	1,385	57.4	615	143.6	2,286	37.5	0.730	243

Critical-difference diagram - p99 latency (ms)

Friedman chi-square = 56.7, $p = 5.39 \times 10^{-9}$ over $N=9$ tasks, $k=10$ stacks. Neményi post-hoc at the 5% level; stacks joined by a bar are not distinguishable.




Figure 3: Critical-difference diagram for p99 latency. Stacks joined by a bar are not distinguishable at the 5% level.

8 Limitations

One machine, one database, loopback networking, one shape of workload. A four-core laptop is not a server, and the database sharing a socket with everything else compresses gaps that a larger machine would widen. The ratios describe how these stacks handle a database-bound transactional API on this hardware; the absolute numbers belong to this machine.

Nothing here speaks to streaming, CPU-bound compute, cold-start behaviour, or performance across a real network.

The ordering *within* a language is not a result. Rust’s hyper, Axum and Actix fall inside the critical difference of one another, as do Go’s Fiber, net/http and Gin; they change places depending on which repetitions are included. The language tiers are robust, the intra-tier ranking is not.

Five repetitions per cell is few. The intervals quoted are percentile bootstrap intervals on a median of five: an honest width, not a precise coverage guarantee. An earlier iteration of this suite with three repetitions and no machine-speed probe produced a 66% median spread across repetitions, and the ranking it implied was partly an artifact of autovacuum and of unpinned neighbouring processes. That failure is the reason the probe exists.

Table 3: Average Friedman ranks over the nine tasks (1 = best).

p99 latency		CPU per request	
stack	rank	stack	rank
go-nethttp	2.56	rust-hyper	1.22
go-fiber	2.67	rust-axum	1.78
rust-axum	3.56	rust-actix	3.22
rust-hyper	4.22	go-fiber	3.89
go-gin	4.44	go-nethttp	5.39
rust-actix	5.00	go-gin	6.39
ts-bun-only	6.00	ts-bun-only	6.44
ts-bun-hono	8.44	ts-bun-hono	7.67
go-gin-gorm	8.89	ts-node-fastify	9.00
ts-node-fastify	9.22	go-gin-gorm	10.00

Table 4: Pre-specified contrasts on saturated throughput. Ratios above one favour the first stack; intervals are bootstrap 95%.

contrast	question	ratio	95% CI	p (Holm)
ts-bun-only vs ts-node-fastify	Bun vs Node, both TypeScript	1.72	[1.51, 1.83]	1.7e-05
go-gin vs go-gin-gorm	Cost of the ORM, router held constant	1.84	[1.54, 2.56]	1.7e-05
ts-bun-only vs ts-bun-hono	Zero-dependency Bun vs Hono + node-postgres	1.33	[1.18, 1.42]	2.2e-05
rust-axum vs ts-bun-only	Best Rust vs best TypeScript	1.45	[1.30, 1.64]	4.7e-05
rust-axum vs go-fiber	Best Rust vs best Go	0.99	[0.86, 1.07]	7.1e-01

9 Future work

A second machine class would separate hardware-specific effects from stack effects. Adding Hono over Bun.sql would disentangle router cost from driver cost in the TypeScript tier, which the present design confounds. Running the database on a separate host would restore the network round-trip that this loopback setup removes.

10 Availability

Source, raw per-run results, analysis scripts and the generator for every figure in this report are at <https://github.com/hallelxl2/polyglot-bench> under the MIT licence. Every number here is produced by `analysis/stats.py` from the run files; none is transcribed by hand.

References

- [1] TechEmpower. *Framework Benchmarks, Round 23*, March 2025. <https://www.techempower.com/benchmarks/>
- [2] J. Demšar. *Statistical Comparisons of Classifiers over Multiple Data Sets*. Journal of Machine Learning Research, 7:1–30, 2006.




Figure 4: The five planned contrasts.

What a request costs you

Down-and-left is cheaper. Median across the three workloads at saturation; Rust holds an order of magnitude less memory than Node for the same work.




Figure 5: CPU and memory cost per request; bubble area is saturated throughput.